

Task Title: Customer Churn Prediction for a Bank

Importing the libraries for the task

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, classification_report,
accuracy_score, jaccard_score, log_loss
import itertools
import seaborn as sns
```

Load the data from the CSV File

```
Data = pd.read_csv('Churn_Modelling.csv')
```

Get data characteristics to better train our model for more accurate results

```
Data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   RowNumber             10000 non-null  int64  
 1   CustomerId            10000 non-null  int64  
 2   Surname               10000 non-null  object  
 3   CreditScore            10000 non-null  int64  
 4   Geography             10000 non-null  object  
 5   Gender                10000 non-null  object  
 6   Age                   10000 non-null  int64  
 7   Tenure                10000 non-null  int64  
 8   Balance                10000 non-null  float64 
 9   NumOfProducts         10000 non-null  int64  
10   HasCrCard             10000 non-null  int64  
11   IsActiveMember        10000 non-null  int64  
12   EstimatedSalary       10000 non-null  float64 
13   Exited                10000 non-null  int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB

Data.duplicated().sum()
```

```
0
```

```
Data.isna().sum()
```

```
RowNumber      0
CustomerId      0
Surname         0
CreditScore    0
Geography      0
Gender         0
Age            0
Tenure         0
Balance        0
NumOfProducts  0
HasCrCard      0
IsActiveMember 0
EstimatedSalary 0
Exited         0
dtype: int64
```

```
Data.dtypes
```

```
RowNumber      int64
CustomerId      int64
Surname        object
CreditScore    int64
Geography      object
Gender         object
Age            int64
Tenure         int64
Balance        float64
NumOfProducts  int64
HasCrCard      int64
IsActiveMember int64
EstimatedSalary float64
Exited         int64
dtype: object
```

I have noticed that the churn column has been labeled as 'Exited'. Now for me that seems a little confusing and misleading for both me and any layman reader so I am going to remove that column.

```
Data['Churn'] = Data['Exited']
Data.drop(['Exited'], axis=1, inplace=True)
```

Lets see if that worked

```
Data.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
\							
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Churn
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

Perfect!

Data.describe()

	RowNumber	CustomerId	CreditScore	Age
Tenure \				
count	10000.000000	1.000000e+04	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800
std	2886.89568	7.193619e+04	96.653299	10.487806
min	1.00000	1.556570e+07	350.000000	18.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000
75%	7500.25000	1.575323e+07	718.000000	44.000000
max	10000.00000	1.581569e+07	850.000000	92.000000

	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
count	10000.000000	10000.000000	10000.000000	10000.000000	
mean	76485.889288	1.530200	0.70550	0.515100	
std	62397.405202	0.581654	0.45584	0.499797	
min	0.000000	1.000000	0.00000	0.000000	
25%	0.000000	1.000000	0.00000	0.000000	
50%	97198.540000	1.000000	1.00000	1.000000	
75%	127644.240000	2.000000	1.00000	1.000000	
max	250898.090000	4.000000	1.00000	1.000000	

	EstimatedSalary	Churn
count	10000.000000	10000.000000
mean	100090.239881	0.203700
std	57510.492818	0.402769
min	11.580000	0.000000
25%	51002.110000	0.000000
50%	100193.915000	0.000000
75%	149388.247500	0.000000
max	199992.480000	1.000000


Now we will visualise our data to better understand it.

```

categorical_features =
Data.select_dtypes(include=[object]).columns[Data.select_dtypes(include=[object]).columns != 'Surname']

plt.figure(figsize = (25,15))
for i, feature in enumerate(categorical_features):
    plt.subplot(3,5,i + 1)
    sns.set_palette('Paired')
    sns.set_style("ticks")
    ax = sns.countplot(x = feature, data = Data)
    ax.set_xticklabels(ax.get_xticklabels(), rotation=45, ha="right")

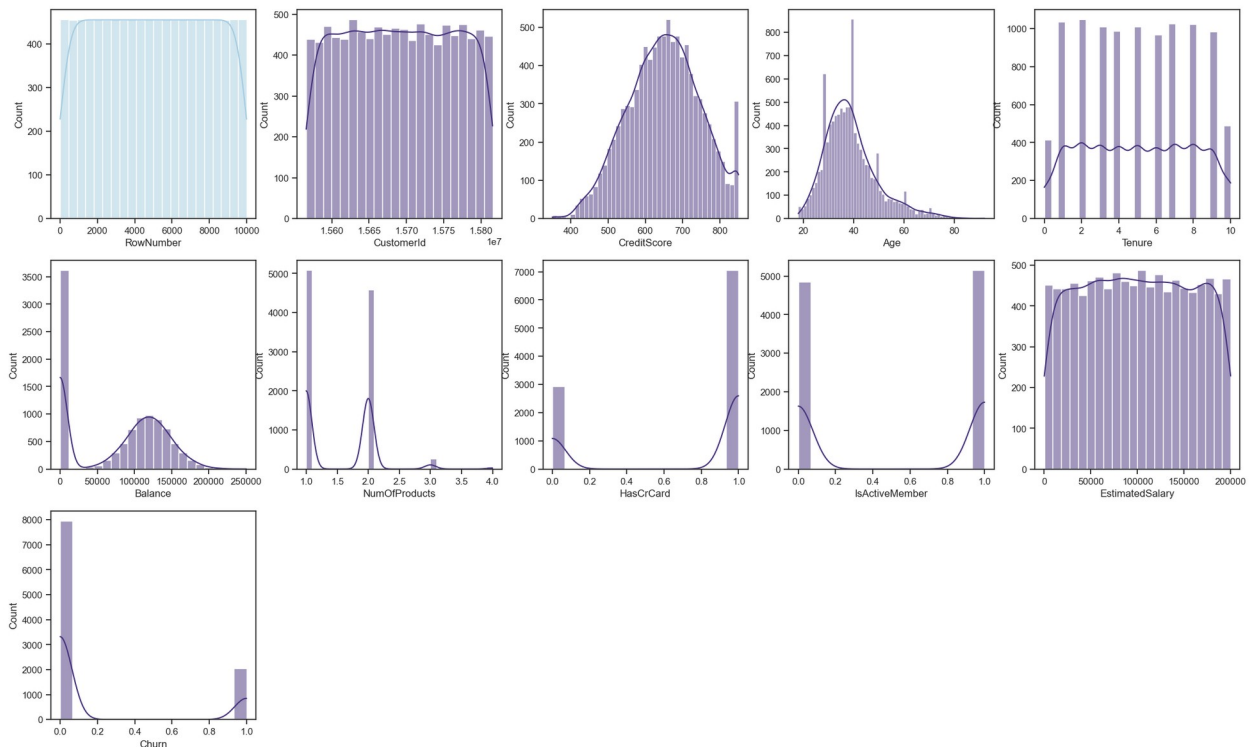
```



```
# Assuming Data is your dataset
numeric_features = Data.select_dtypes(include=[np.number])

# Calculate the number of rows and columns for subplots
num_features = numeric_features.shape[1]
num_rows = (num_features - 1) // 5 + 1
num_cols = min(num_features, 5)

plt.figure(figsize=(25, 5 * num_rows))
for i, feature in enumerate(numeric_features.columns):
    plt.subplot(num_rows, num_cols, i + 1)
    sns.set(palette='viridis')
    sns.set_style("ticks")
    sns.histplot(Data[feature], kde=True)
    plt.xlabel(feature)
    plt.ylabel("Count")
```



For further operations, such as data normalization etc., we need to convert the data into a numpy array.

```
# Select all int64 columns except for the 'Churn' column
cols =
Data.select_dtypes(include=['int64']).columns[Data.select_dtypes(include=['int64']).columns != 'Churn']
# Create a new DataFrame with the selected columns
X = np.asarray(Data[cols])
# Display the first 5 rows of the new DataFrame
X[0:5]
```

```
array([[ 1, 15634602, 619, 42, 2, 1,
        1, 1],
       [ 2, 15647311, 608, 41, 1, 1,
        0, 1],
       [ 3, 15619304, 502, 42, 8, 3,
        1, 0],
       [ 4, 15701354, 699, 39, 1, 2,
        0, 0],
       [ 5, 15737888, 850, 43, 2, 1,
        1, 1]], dtype=int64)
```

```
y = np.asarray(Data['Churn'])
y [0:5]
```

```
array([1, 0, 1, 0, 0], dtype=int64)
```

```

X = preprocessing.StandardScaler().fit(X).transform(X)
X[0:5]

array([[ -1.73187761, -0.78321342, -0.32622142,  0.29351742, -
 1.04175968,
        -0.91158349,  0.64609167,  0.97024255],
       [ -1.7315312 , -0.60653412, -0.44003595,  0.19816383, -
 1.38753759,
        -0.91158349, -1.54776799,  0.97024255],
       [ -1.73118479, -0.99588476, -1.53679418,  0.29351742, -
 1.03290776,
        2.52705662,  0.64609167, -1.03067011],
       [ -1.73083838,  0.14476652,  0.50152063,  0.00745665, -
 1.38753759,
        0.80773656, -1.54776799, -1.03067011],
       [ -1.73049197,  0.65265871,  2.06388377,  0.38887101, -
 1.04175968,
        -0.91158349,  0.64609167,  0.97024255]])

X_train, X_test, y_train, y_test =
train_test_split(X,y,test_size=0.2,random_state=4)
print ('Train set:', X_train.shape,  y_train.shape)
print ('Test set:', X_test.shape,  y_test.shape)

Train set: (8000, 8) (8000,)
Test set: (2000, 8) (2000,)

```

So the data has a 80/20 split now. 80% of the data is in the training set and 20% of the data is in the test set. Now we will create a logistic regression model and train it on the training set. We will name our model 'LogRes'.

```

LogRes = LogisticRegression(C = 0.1,
solver='liblinear').fit(X_train,y_train)
LogRes

LogisticRegression(C=0.1, solver='liblinear')

yhat = LogRes.predict(X_test)
yhat

array([0, 0, 0, ..., 0, 0, 0], dtype=int64)

```

Now we have obtained the data into fine groups but we also need to find the probability of each possibility of the data. For this we will use the function below.

```

yhat_prob = LogRes.predict_proba(X_test)
yhat_prob

array([[0.89185868, 0.10814132],
       [0.7369563 , 0.2630437 ]],

```

```

        [0.92860437, 0.07139563],
        ...,
        [0.73854618, 0.26145382],
        [0.82689892, 0.17310108],
        [0.86627779, 0.13372221]])

jaccard_score(y_test, yhat, pos_label=0)

0.8066666666666666

```

We have got a 80% accuracy on the test set. This is a good accuracy. However, it can be better.

```

def Confusion_Matrix(ConfMat, Categories, normalize=False, title =
'Churn Prediction Confusion Matrix', cmap = plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    if normalize:
        ConfMat = ConfMat.astype('float') / ConfMat.sum(axis=1)[:,
np.newaxis]
        print("Normalized Confusion Matrix")
    else:
        print('Confusion Matrix Without Normalization')
    print(ConfMat)

    plt.imshow(ConfMat, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(Categories))
    plt.xticks(tick_marks, Categories, rotation=45)
    plt.yticks(tick_marks, Categories)

    fmt = '.2f' if normalize else 'd'
    threshold = ConfMat.max() / 2.
    for i, j in itertools.product(range(ConfMat.shape[0]),
range(ConfMat.shape[1])):
        plt.text(j, i, format(ConfMat[i, j], fmt),
                horizontalalignment="center",
                color="white" if ConfMat[i, j] > threshold else
"black")

    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.tight_layout()

# Path: Customer Churn Prediction.ipynb
print(confusion_matrix(y_test, yhat, labels=[1,0]))

```



```

[[ 50 346]
 [ 31 1573]]

Conf_Matrix = confusion_matrix(y_test, yhat, labels=[1,0])
print(Conf_Matrix)

[[ 50 346]
 [ 31 1573]]

np.set_printoptions(precision=3, suppress=True)

```

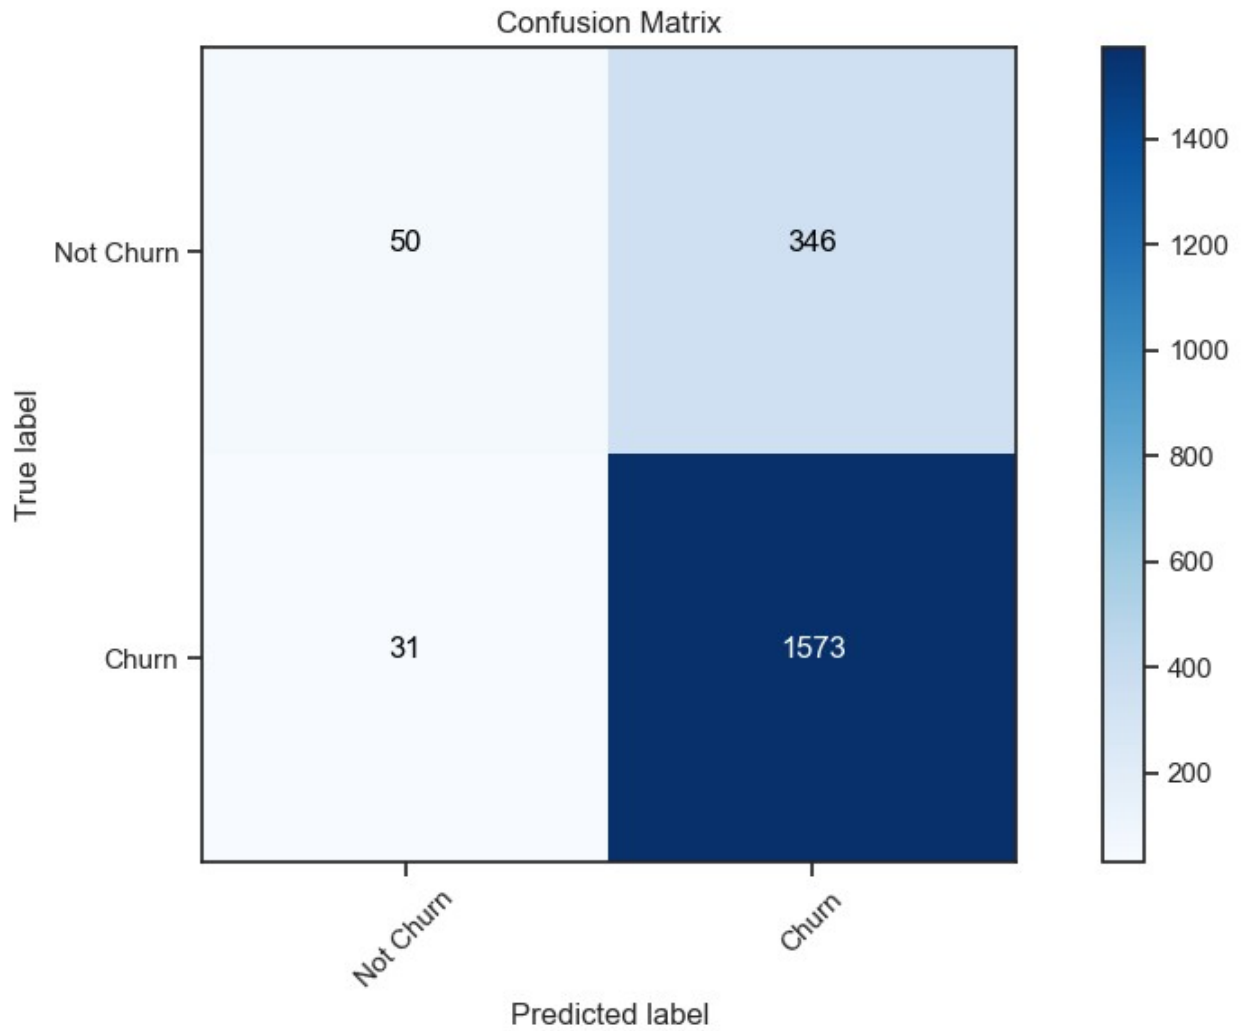
We will now obtain a Confusion Matrix below which is a classification of our predictions into a table. The columns represent the predicted labels and the rows represent the actual labels. The cells represent the number of predictions made by the model. The diagonal cells represent the number of correct predictions made by the model. The off-diagonal cells represent the number of incorrect predictions made by the model. The Confusion Matrix is a great way to visualize the performance of our model.

```

plt.figure(figsize=(10, 6))
Confusion_Matrix(Conf_Matrix, Categories= ['Not Churn',
'Churn'], normalize= False , title='Confusion Matrix', cmap='Blues')

Confusion Matrix Without Normalization
[[ 50 346]
 [ 31 1573]]

```



```
print(classification_report(y_test, yhat))
```

	precision	recall	f1-score	support
0	0.82	0.98	0.89	1604
1	0.62	0.13	0.21	396
accuracy			0.81	2000
macro avg	0.72	0.55	0.55	2000
weighted avg	0.78	0.81	0.76	2000

```
log_loss(y_test, yhat_prob)
```

```
0.4386196611392611
```

```
accuracy_score(y_test, yhat)
```

```
0.8115
```

A slight improvement from our previous result of 80%. Lets try another algorithm.

Algorithm: Random Forest

Now, as per our instructions, we also have the liberty to run the Random Forest algorithm on our data.

```
import pandas as pd
from sklearn import metrics
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegressionCV
from imblearn.over_sampling import SMOTE
from sklearn.metrics import confusion_matrix, auc, roc_curve,
precision_score, recall_score

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Perform feature scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

As you can see, the starting steps of libraries importing and train/test data splitting are the same. The only difference that you might notice is that we have imported the 'RandomForestClassifier' unlike the previous model in which we imported 'LogisticRegression' from the linear_model library of sklearn.

```
rf_classifier = RandomForestClassifier(random_state=42)
```

We will now test our models accuracy.

```
rf_classifier.fit(X_train_scaled, y_train)

# Make predictions on the test data
y_pred = rf_classifier.predict(X_test_scaled)

# Calculate the train accuracy of the classifier
tr_accuracy = rf_classifier.score(X_train_scaled, y_train)

# Calculate the accuracy of the classifier
ts_accuracy = accuracy_score(y_test, y_pred)

print("Train Accuracy:", f'{round(tr_accuracy, 4) * 100}%')
print("Test Accuracy:", f'{round(ts_accuracy, 4) * 100}%')
```

Train Accuracy: 100.0%
Test Accuracy: 86.15%

Already, the model is much better than the final result we got from the logistic regression model. This seems promising.

```
c_matrix = confusion_matrix(y_test, y_pred)

c_matrix_display = metrics.ConfusionMatrixDisplay(confusion_matrix =
c_matrix)


c_matrix_display.plot()
print("Confusion Matrix:")
plt.show()

from sklearn.metrics import classification_report
print("Classification Report:\n", classification_report(y_test,
y_pred))

yhat_prob = rf_classifier.predict_proba(X_test)
yhat_prob

print("Log Loss: ", log_loss(y_test, yhat_prob))

Confusion Matrix:
```



Classification Report:				
	precision	recall	f1-score	support
0	0.88	0.96	0.92	1607
1	0.75	0.44	0.56	393
accuracy			0.86	2000
macro avg	0.81	0.70	0.74	2000
weighted avg	0.85	0.86	0.85	2000
Log Loss: 0.41215324172510215				

SMOTED Version of Algorithm: Random Forest

Now, a way that we can modify our Random Forest model is to using the function SMOTE. It is used to perform a balance class distribution on an otherwise imbalanced dataset. It is an oversampling method that creates synthetic samples from the minor class. This is done by selecting similar records and altering that records one by one to create a new synthetic record. This is done until the number of records in both classes is equal.

```
oversample = SMOTE(k_neighbors=55, random_state=42)
X_features_smoted, y_target_smoted = oversample.fit_resample(X, y)
y_target_smoted = pd.Series(y_target_smoted)
```

```

X_train, X_test, y_train, y_test = train_test_split(X_features_smoted,
y_target_smoted, test_size=0.2, random_state=42)

# Perform feature scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

rf_classifier.fit(X_train_scaled, y_train)

# Make predictions on the test data
y_pred = rf_classifier.predict(X_test_scaled)

# Calculate the train accuracy of the classifier
tr_accuracy = rf_classifier.score(X_train_scaled, y_train)

# Calculate the accuracy of the classifier
ts_accuracy = accuracy_score(y_test, y_pred)

print("Train Accuracy:", f'{round(tr_accuracy, 4) * 100}%')
print("Test Accuracy:", f'{round(ts_accuracy, 4) * 100}%')

Train Accuracy: 100.0%
Test Accuracy: 86.91%

```

We have increased our accuracy from 85% to almost 87% after using the SMOTE technique.

```

c_matrix = confusion_matrix(y_test, y_pred)

c_matrix_display = metrics.ConfusionMatrixDisplay(confusion_matrix =
c_matrix)

c_matrix_display.plot()
print("Confusion Matrix:")
plt.show()


from sklearn.metrics import classification_report
print("Classification Report:\n", classification_report(y_test,
y_pred))

yhat_prob = LogRes.predict_proba(X_test)
yhat_prob

print("Log Loss: ", log_loss(y_test, yhat_prob))

Confusion Matrix:

```



Classification Report:				
	precision	recall	f1-score	support
0	0.86	0.89	0.87	1633
1	0.88	0.85	0.86	1553
accuracy			0.87	3186
macro avg	0.87	0.87	0.87	3186
weighted avg	0.87	0.87	0.87	3186
Log Loss: 0.7666645884787358				

We are satisfied that this is the highest accuracy that can be achieved by this method. We will now try to improve the accuracy by using other methods.

Algorithm: Gradient Boosting

A final algorithm that we would use is the Gradient Boosting.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
```

```

from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import confusion_matrix , classification_report ,
accuracy_score, jaccard_score, log_loss
import itertools
import seaborn as sns

X_train, X_test, y_train, y_test =
train_test_split(X,y,test_size=0.2,random_state=4)
print ('Train set:', X_train.shape,  y_train.shape)
print ('Test set:', X_test.shape,  y_test.shape)

Train set: (8000, 8) (8000,)
Test set: (2000, 8) (2000,)

```

We will create a model, same as before.

```

GradBoost =
GradientBoostingClassifier(n_estimators=100).fit(X_train,y_train)
GradBoost

GradientBoostingClassifier()

yhat = GradBoost.predict(X_test)
yhat

array([0, 0, 0, ..., 0, 1, 0], dtype=int64)

yhat_prob = GradBoost.predict_proba(X_test)
yhat_prob

array([[0.89 , 0.11 ],
       [0.916, 0.084],
       [0.968, 0.032],
       ...,
       [0.804, 0.196],
       [0.023, 0.977],
       [0.954, 0.046]])

jaccard_score(y_test, yhat, pos_label=0)

0.8526663001649258

```

Our jaccard score shows that our model is 84% accurate/similar to the actual data. This is a good score, but we can do better. We will try to improve our model.

```

def Confusion_Matrix(ConfMat, Categories, normalize=False, title =
'Churn Prediction Confusion Matrix', cmap = plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.

```



```

"""
if normalize:
    ConfMat = ConfMat.astype('float') / ConfMat.sum(axis=1)[:,
np.newaxis]
    print("Normalized Confusion Matrix")
else:
    print('Confusion Matrix Without Normalization')
print(ConfMat)

plt.imshow(ConfMat, interpolation='nearest', cmap=cmap)
plt.title(title)
plt.colorbar()
tick_marks = np.arange(len(Categories))
plt.xticks(tick_marks, Categories, rotation=45)
plt.yticks(tick_marks, Categories)

fmt = '.2f' if normalize else 'd'
threshold = ConfMat.max() / 2.
for i, j in itertools.product(range(ConfMat.shape[0]),
range(ConfMat.shape[1])):
    plt.text(j, i, format(ConfMat[i, j], fmt),
             horizontalalignment="center",
             color="white" if ConfMat[i, j] > threshold else
"black")

plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.tight_layout()

# Path: Customer Churn Prediction.ipynb
print(confusion_matrix(y_test, yhat, labels=[1,0]))

[[ 181  215]
 [  53 1551]]

Conf_Matrix = confusion_matrix(y_test, yhat, labels=[1,0])
print(Conf_Matrix)


[[ 181  215]
 [  53 1551]]

np.set_printoptions(precision=3, suppress=True)

plt.figure(figsize=(10, 6))
Confusion_Matrix(Conf_Matrix, Categories= ['Not Churn',
'Churn'],normalize= False , title='Confusion Matrix', cmap='Blues')

Confusion Matrix Without Normalization
[[ 181  215]
 [  53 1551]]

```



```
print(classification_report(y_test,yhat))
```

	precision	recall	f1-score	support
0	0.88	0.97	0.92	1604
1	0.77	0.46	0.57	396
accuracy			0.87	2000
macro avg	0.83	0.71	0.75	2000
weighted avg	0.86	0.87	0.85	2000

```
log_loss(y_test,yhat_prob)
```

```
0.34300451502713625
```

So we have achieved about 87% accuracy from our Gradient Boost Classifier model.

Algorithm Comparision:

We will now compare our Algorithms to see which one gave us the best answer.

```
Algo_List = [LogRes, rf_classifier, GradBoost]

X_train, X_test, y_train, y_test = train_test_split(X_features_smoted,
y_target_smoted, test_size=0.2, random_state=42)

# Perform feature scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

Algo_columns = []
Algo_compare = pd.DataFrame(columns = Algo_columns)

row_index = 0

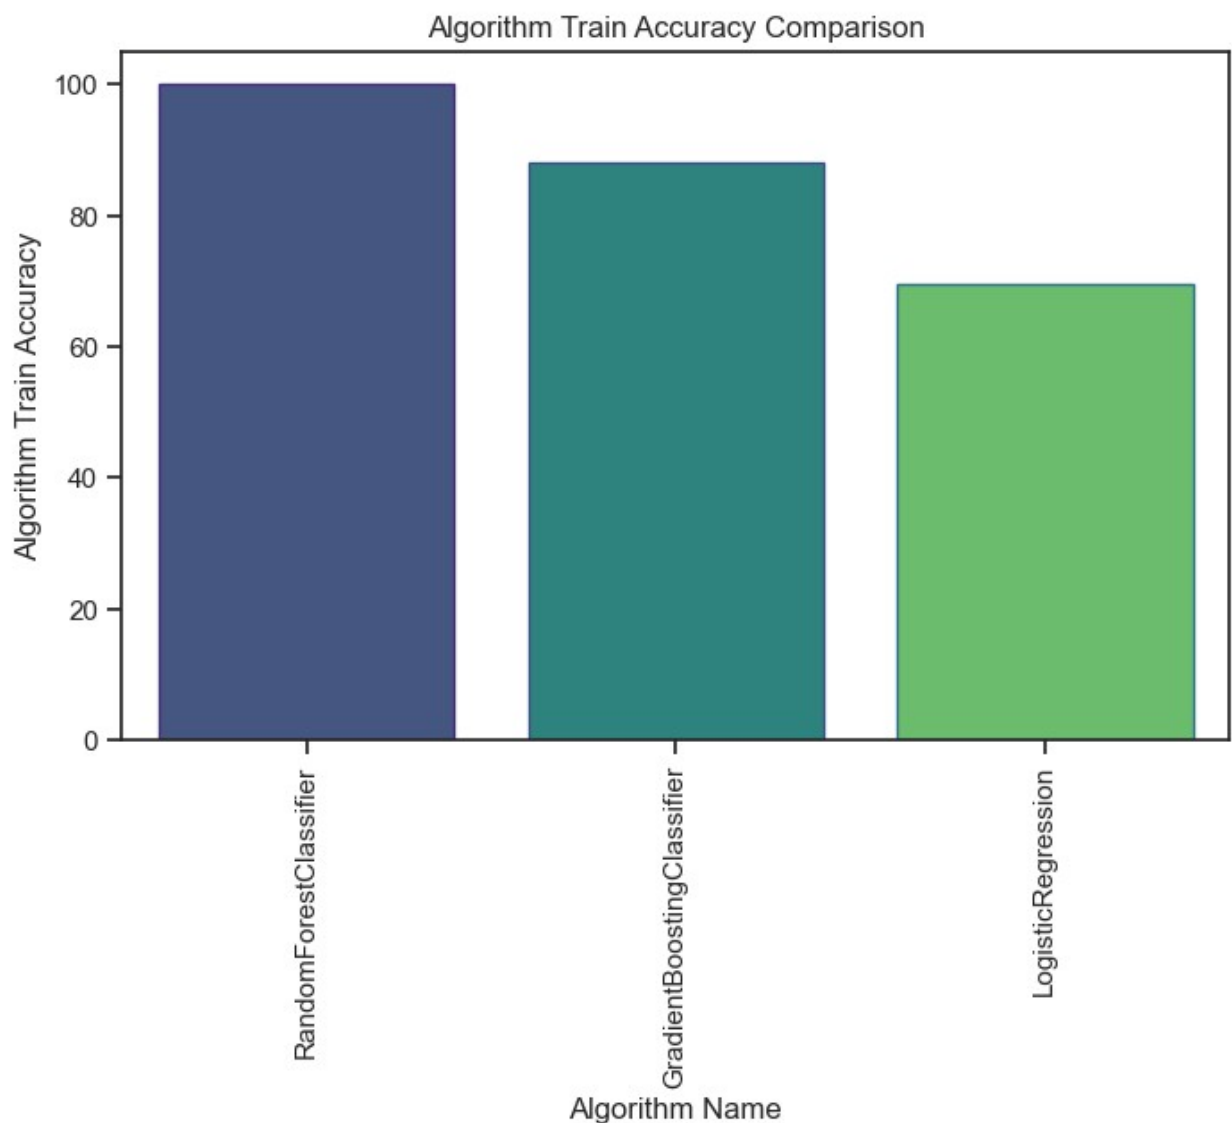
for alg in Algo_List:
    predicted = alg.fit(X_train_scaled,
y_train).predict(X_test_scaled)
    fp, tp, th = roc_curve(y_test, predicted)
    MLA_name = alg.__class__.__name__
    Algo_compare.loc[row_index, 'Algorithm Name'] = MLA_name
    Algo_compare.loc[row_index, 'Algorithm Train Accuracy'] =
round(alg.score(X_train_scaled, y_train), 4) * 100
    Algo_compare.loc[row_index, 'Algorithm Test Accuracy'] =
round(alg.score(X_test_scaled, y_test), 4) * 100
    Algo_compare.loc[row_index, 'Algorithm Precision'] =
round(precision_score(y_test, predicted), 4) * 100
    Algo_compare.loc[row_index, 'Algorithm Recall'] =
round(recall_score(y_test, predicted), 4) * 100
    Algo_compare.loc[row_index, 'Algorithm AUC'] = round(auc(fp, tp),
4) * 100
    row_index+=1

Algo_compare.sort_values(by = ['Algorithm Train Accuracy'], ascending
= False, inplace = True)
Algo_compare
```

	Algorithm Name	Algorithm Train Accuracy \		
1	RandomForestClassifier	100.00		
2	GradientBoostingClassifier	88.01		
0	LogisticRegression	69.69		
	Algorithm Test Accuracy	Algorithm Precision	Algorithm Recall \	
1	86.91	87.82	84.93	
2	88.04	89.12	85.96	
0	69.55	68.65	69.09	

	Algorithm AUC
1	86.86
2	87.99
0	69.54

```
plt.subplots(figsize=(8,5))
sns.barplot(x="Algorithm Name", y="Algorithm Train Accuracy", data=Algo_compare, palette='viridis', edgecolor=sns.color_palette('viridis',7))
plt.xticks(rotation=90)
plt.title('Algorithm Train Accuracy Comparison')
plt.show()
```



Based on our results, we are able to affirm that despite RandomForestClassifier having the best training accuracy out of all three, it has become overfitted and only has 86% test accuracy.

Therefore, based on training and test data accuracy and the confusion matrix, we can conclude that Gradient Boost is the best model for this dataset.